

Affix the authentication label over this field.

Surname:

First name:

Matr.-No.:

**Institut für Mess- und Regelungstechnik
mit Maschinenlaboratorium
Karlsruher Institut für Technologie (KIT)
Prof. Dr.-Ing. C. Stiller**

Exam in “Machine Vision”

Date of exam: March 18, 2024

Time of exam: 8:00-9:00

Question 1

(5 points)



Different kernels are convoluted with the graylevel image above. Below you can see the output of the convolution operation. Which kernel belongs to which output image? Draw a line between the matching output image and kernel.

$$\begin{bmatrix} 0 & 2 & 0 \\ 2 & 0 & 2 \\ 0 & 2 & 0 \end{bmatrix}$$


$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & 2 & -1 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$





Question 2

(8 points)

Implement a Python function `is_underexposed` that evaluates whether a video is underexposed.

A video is considered underexposed if a certain percentage of all pixels within the video is underexposed.

The function takes three arguments:

- A video of shape $[F, C, H, W]$, with F being the amount of frames, C the channels, H the height, and W the width.
- A threshold defining the percentage of underexposed pixels from which on a video is considered underexposed.
- A threshold defining whether a pixel is underexposed.

```
import numpy as np
```

```
def is_underexposed(video, t_vid = 0.3, t_pixel = 0):  
    # convert RGB video to grayscale video  
    gray_video = to_grayscale(video)  
  
    # @student: implement function
```



Question 3

(6 points)

The figure below shows a graylevel picture with 5 rows and 5 columns.

- (a) Apply the connected components labeling algorithm (CCL) to segment the image. Two gray values should be treated as similar only if they are equal, i.e. gray pixels and white pixels in the figure below are always dissimilar. Visualize the processing of CCL by inserting into each cell in the figure one of the following letters to indicate the processing step:


N if CCL initializes a new segment at this pixel

M if CCL merges two segments when passing this pixel


L if CCL assigns the pixel to the segment of its left neighbor, but not to the segment of its upper neighbor

U if CCL assigns the pixel to the segment of its upper neighbor, but not to the segment of its left neighbor

B if CCL assigns the pixel to the segment of its left neighbor and the left and upper neighbor already belong to the same segment



- (b) How many segments are finally found?



Question 4

(8 points)

In a project you need to calibrate a camera. Your colleague already executed two different camera calibration methods resulting in two sets of calibration parameters R_1 , t_1 , A_1 and R_2 , t_2 , A_2 .

$$R_1 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$R_2 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$t_1 = \begin{bmatrix} 0.5 \\ 0.5 \\ 0 \end{bmatrix}$$

$$t_2 = \begin{bmatrix} 0.5 \\ 0.5 \\ 0 \end{bmatrix}$$


$$A_1 = \begin{bmatrix} 20 & 0 & 1000 \\ 0 & 40 & 1000 \\ 0 & 0 & 1 \end{bmatrix}$$

$$A_2 = \begin{bmatrix} 20 & 0 & 1000 \\ 0 & 30 & 1000 \\ 0 & 0 & 1 \end{bmatrix}$$

Furthermore, he recorded one world and image point pair:

$$\begin{bmatrix} 1.0 \\ 0.5 \\ 0.5 \end{bmatrix} \Rightarrow \begin{bmatrix} 1060 \\ 1065 \end{bmatrix}$$

Unfortunately, due to his child being sick, he cannot attend work this week and you need to decide which calibration parameters to use. Calculate and explain your reasoning.




Question 5

(8 points)

Assume that we train four different classifiers A , B , C , D for a certain classification task and test them on a test set. The figure below shows a Venn diagram of the number of misclassified test examples. Each ellipse represents one of the four classifiers. The numbers provide how many test examples are misclassified by the respective classifiers. E.g., there are 15 examples misclassified by classifier A while all the other classifiers are correct. There are 5 test examples which are misclassified by classifier A and C while the classifiers B and D are correct. There is one test example that is misclassified by the classifiers A , C , and D , while classifier B is correct.

For each possible ensemble of three of these classifiers calculate the number of test examples that are misclassified.




ensemble members	misclassified test examples
A, B, C	
A, B, D	
A, C, D	
B, C, D	



Question 6

(3+3+3 points)



Assume the architecture of the neural network in the image above predicts four different classes of images. The input of the architecture is an image of shape $[C, H, W]$ with $C = 3$ (channels), $H = 12$ (height), and $W = 12$ (width).

- (a) Assume a filters dimension is $[3, 3, 3]$ and 10 filters are learned in the convolution layer. Calculate how many parameters the convolution layer has.
- (b) After the convolution layer is applied to the image, you end up with a feature mask of shape $[10, 10, 10]$. To predict the class probabilities, a linear layer is applied to the flattened feature masks. Calculate how many parameters the linear layer has.
- (c) Calculate how many multiplications are required to execute the convolution layer.



Question 7


(1+2+3 points)

Maria wants to train a deep-learning model to classify images of eight different road signs. She collects a data set of 500 images. She uses 400 images for training, 50 images for validation, and 50 images for testing.

To select a model she looks at a leaderboard and chooses the best performing model available. The model is a transformer model and has over one billion parameters. For training, she shuffles the training data and trains for 50 epochs.


After training, she achieves a perfect accuracy of 100% on the training data. However, on the validation data she only achieves an accuracy of 12.4%.

- (a) Name the phenomenon that Maria is faced with.
- (b) Explain why this occurs.
- (c) Name three countermeasures to avoid this phenomenon.



Question 8

(10 points)



In the image above you can see 2D points extracted by a machine vision algorithm which represent a circle. However, to further process this circle your application requires it in a parametric form.

Extend the RANSAC algorithm to detect circles in an unordered set of 2D points. Based on the RANSAC algorithm, describe your extension step by step.

Explain your steps mathematically. You do not need to fit a circle to the given sample data.

Hint: the equation of a circle is $(x - x_c)^2 + (y - y_c)^2 = r^2$

